

Statement of Interest

“Predicting the next geometric element, not the next word”

I am writing to express my strong interest in the research project titled "Large Language Model for Computing Midcurve of a Thin Polygon." The concept of generating midcurves for 2D geometric profiles using Large Language Models (LLMs) presents an exciting and challenging problem with implications that extend beyond the domain of machine design. As an AI researcher with a focus on natural language processing and machine learning, I am deeply motivated to contribute to this project and explore the boundaries of what LLMs can achieve.

The problem at hand, namely, finding the midcurve of a 2D geometric profile, is a complex and longstanding challenge in the field of computer-aided design. Despite various attempts with traditional geometric techniques like Medial Axis Transform, Chordal Axis Transform, Thinning, and Pairing, the problem remains unsolved due to the diversity and complexity of 2D shapes and their intricate connections. The proposed research project, MidcurveNN, seeks to evaluate whether Neural Networks, and particularly Large Language Models, can provide a solution to this problem.

The problem statement focuses on the goal of generating the midcurve for a closed polygon, which can be represented as a set of connected lines or points. This problem essentially translates to a graph summarization and dimension-reduction task, analogous to text summarization, where the challenge lies in preserving the underlying structure while reducing complexity. It poses several critical questions about how geometric data can be represented, transformed, and summarized in a way that can be fed into neural networks effectively.

One of the key challenges is that geometric shapes are not always easily representable as sequences, unlike language data. Shapes like 'L', 'I', and 'T' cannot be directly modeled as sequences of points or lines, making it a non-trivial task for conventional sequence-to-sequence networks. The research aims to explore innovative approaches to represent and convert 2D linear profile shapes into vector forms suitable for neural network processing. This involves addressing the issue of how to incorporate geometric information into graph structures, which inherently possess topological, not geometrical, representations.

The proposed approach of utilizing Large Language Models, such as GPT-4, to generate midcurves by providing prompts is a novel and promising direction. By formulating the problem as a sequence generation task for LLMs, the research explores the potential of these models to understand complex geometric shapes and produce meaningful midcurve representations. The project's findings demonstrate the current limitations of LLMs in handling geometric data while also highlighting their potential for further development.

The research's real-world applications extend beyond the field of machine design. If successful, the ability of LLMs to predict complex geometric structures opens up possibilities in various domains, such as social network analysis, pharmaceutical research, and more, where graph-based data representations are common.

In conclusion, the research project "Large Language Model for Computing Midcurve of a Thin Polygon" aligns with my research interests and expertise in natural language processing and machine learning. I am eager to contribute my skills and knowledge to address the

challenges presented by this project. The potential to expand the capabilities of Large Language Models beyond text and into the realm of complex geometric data is a compelling avenue for exploration, and I am excited to be a part of this research endeavor.

Large Language Model for Computing Midcurve of a Thin Polygon


ABSTRACT

This paper explores the potential of LLMs (Large Language Models) in computing the midcurve of a 2D geometric profile. The midcurve, a curve equidistant from the bounding curve of the profile, simplifies the representation of the profile-shape while retaining essential geometric information. Despite various approaches such as Medial Axis Transform, Chordal Axis Transform, Thinning and Pairing, the problem remains unsolved due to complexity of shapes and variety of connections. The paper evaluates if the LLMs can be used to generate the midcurve of a 2D geometric profile.

Keywords: Midcurve; Large Language Models; Dimension Reduction; Artificial Intelligence


1. Introduction

A midcurve of a 2D geometric profile is a curve that lies equidistant from the bounding curves of the profile. It is a curve that represents the “middle” of the profile. The midcurve can be used to represent the shape of the profile in a simpler way than the original profile, while still retaining important geometric information about the profile.



Example of Midcurves of 2D Geometric profiles ([ref](#))

Although many approaches like Medial Axis Transform, Chordal Axis Transform, Thinning, Pairing, etc., have been tried for decades, the problem remains unsolved due to complexity of shapes and variety of connections.



Midcurve Approaches ([ref](#))

Research project [MidcurveNN](#) attempts to evaluate if Neural Networks can be used to generate the midcurve of a 2D geometric profile.

2. Problem Statement

- **Goal:** Given a 2D closed shape (closed polygon) find its midcurve (polyline, closed or open)


- **Input:** set of points or set of connected lines, non-intersecting, simple, convex, closed polygon
 - **Output:** another set of points or set of connected lines, open/branched polygons possible
- Essentially, if we consider vertices as nodes and lines as arcs then the polygon/polyline profile is nothing but a Graph and thus the midcurve generation is a Graph Summarization/Dimension-Reduction/Compression issue, i.e. reducing a large graph to a smaller graph preserving its underlying structure, similar to text summarization, which attempts to keep the essence.)

Geometry Dimension-Reduction or Compression can be viewed as encoder-decoder problem where larger input is fed to the encoder side whereas reduced output is generated by the decoder side.

To be able to use Neural Network based encoder decoders, both input and output needs to be in a fixed size vector form. So, can variable shape polygons/polylines be fed to the neural networks? Issues are:


- Shapes cannot be modeled as sequences. Although polygon shape L may appear as sequence of points, it is not.
- All shapes cannot be drawn without lifting a pencil, which is possible for sequences. Say, Shapes like Y or concentric O, cannot be modeled as sequences. So, Midcurve transformation cannot be modeled as Sequence 2 Sequence network.
- How to represent a geometric figure to feed to any Machine/Deep Learning problem, as they need numeric data in vector form?
- How to convert geometric shapes (even after restricting the domain to 2D linear profile shapes) to vectors?
- The closest data structure is graph, but that's predominantly topological and not geometrical. Meaning, graphs represent only connectivity and not spatial positions. Here, nodes would have coordinates and arcs would have curved-linear shape. So, even Graph Neural Networks, which convolute neighbors around a node and pool the output to generate vectors, are not suitable.
- Need RnD to come up with a way to generate geometric graph embedding that will convolute at nodes (having coordinates) around arcs (having geometry, say, a set of point coordinates), then pool them to form a meaningful representation. The real crux would be how to formulate pooling to aggregate all incident curves info plus node coordinates info into a single number!!

So, though graphs are not perfect but with insertion of geometric information they can represent the 2D profiles well. Can Encoder-Decoder architectures (Transformers with Attention) be applied to generate the midcurve?




Midcurve by Encoder Decoder ([ref](#))

Such Transformers taking variable length input and output graphs have not been established. So, would taking this problem to image domain help?



Auto Encoder on images ([ref](#))

Various Image2Image networks have been tried by MidcurveNN project but with limited success. Here is a sample:



Simple Encoder Decoder network in Tensorflow/Keras ([ref](#))

Now evaluating if LLMs (Large Language Models), like can be employed to generate the midcurve.

3. Proposed Idea

The idea is to give prompt telling what needs to be done along with some examples, and see if LLMs can generate shape for the test example. A few shots prompt developed:

You are a geometric transformation program that transforms input 2D polygonal profile to output 1D polyline profile.

Input 2D polygonal profile is defined by set of connected lines with the format as:

input : [line_1, line_2, line_3,...] where lines are defined by two points, where each point is defined by x and y coordinates. So

line_1 is defined as ((x_1, y_1), (x_2,y_2)) and similarly the other lines.

Output is also defined similar to the input as a set of connected lines where lines are defined by two points, where each point is defined by x and y coordinates. So,

output : [line_1, line_2, line_3,...]

Below are some example transformations, specified as pairs of 'input' and the corresponding 'output'. After learning from these examples, predict the 'output' of the last 'input' specified. Do not write code or explain the logic but just give the list of lines with point coordinates as specified for the 'output' format.

input:[((5.0,5.0), (10.0,5.0)), ((10.0,5.0), (10.0,30.0)), ((10.0,30.0), (35.0,30.0)), ((35.0,30.0), (35.0, 35.0)), ((35.0, 35.0), (5.0,35.0)), ((5.0,35.0), (5.0,5.0))]

output: $[((7.5, 5.0), (7.5, 32.5)), ((7.5, 32.5), (35.0, 32.5)), ((35.0, 32.5), (7.5, 32.5))]$

input: $[((5, 5), (10, 5)), ((10, 5), (10, 20)), ((10, 20), (5, 20)), ((5, 20), (5, 5))]$

output: $[((7.5, 5), (7.5, 20))]$


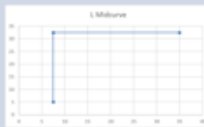
input: $[((0, 25.0), (10.0, 25.0)), ((10.0, 25.0), (10.0, 45.0)), ((10.0, 45.0), (15.0, 45.0)), ((15.0, 45.0), (15.0, 25.0)), ((15.0, 25.0), (25.0, 25.0)), ((25.0, 25.0), (25.0, 20.0)), ((25.0, 20.0), (15.0, 20.0)), ((15.0, 20.0), (15.0, 0)), ((15.0, 0), (10.0, 0)), ((10.0, 0), (10.0, 20.0)), ((10.0, 20.0), (0, 20.0)), ((0, 20.0), (0, 25.0))]$

output: $[((12.5, 0), (12.5, 22.5)), ((12.5, 22.5), (12.5, 45.0)), ((12.5, 22.5), (0, 22.5)), ((12.5, 22.5), (25.0, 22.5))]$



input: $[((0, 25.0), (25.0, 25.0)), ((25.0, 25.0), (25.0, 20.0)), ((25.0, 20.0), (15.0, 20.0)), ((15.0, 20.0), (15.0, 0)), ((15.0, 0), (10.0, 0)), ((10.0, 0), (10.0, 20.0)), ((10.0, 20.0), (0, 20.0)), ((0, 20.0), (0, 25.0))]$

output:

The first input example above represents 'L' shape (shown below) and the second is an 'I', whereas the 3rd is a 'Plus' sign shape.

Profile Data		Profile Picture	Midcurve Data		Midcurve Picture
5.0	5.0		7.5	5.0	
10.0	5.0		7.5	32.5	
10.0	30.0		35.0	32.5	
35.0	30.0		7.5	32.5	
35.0	35.0				
5.0	35.0				

The last shape for which LLM has been asked for the answer is actually a 'T' shape. The picture below shows the correct/actual answer as well.

Profile Data		Profile Picture	Midcurve Data		Midcurve Picture
0	25.0		12.5	0	
25.0	25.0		12.5	22.5	
25.0	20.0		25.0	22.5	
15.0	20.0		0	22.5	
15.0	0				
10.0	0				
10.0	20.0				
0	20.0				


And the outputs computed by various LLMs (ChatGPT, Perplexity AI, Bard) , along with the real/actual answer:

Actual: $[((12.5, 0), (12.5, 22.5)), ((12.5, 22.5), (25.0, 22.5)), ((12.5, 22.5), (0, 22.5))]$


ChatGPT: $[((2.5, 0), (2.5, 22.5)), ((2.5, 22.5), (2.5, 45.0)), ((2.5, 22.5), (25.0, 22.5)), ((2.5, 22.5), (12.5, 22.5)), ((2.5, 22.5), (0, 22.5)), ((2.5, 22.5), (25.0, 22.5))]$
 Perplexity: $[((12.5, 0), (12.5, 22.5)), ((12.5, 22.5), (12.5, 45.0)), ((12.5, 22.5), (0, 22.5)), ((12.5, 22.5), (25.0, 22.5))]$
 Bard: $[((12.5, 0), (12.5, 25.0)), ((12.5, 25.0), (25.0, 25.0)), ((25.0, 25.0), (25.0, 0))]$

Visually here is how results from different LLMs look:


Actual/Expected




ChatGPT



Perplexity AI



Bard



All of the above have failed. Even latest (Oct 2023), the results are:

- llama 7B g4_0 ggml: $(8, 17)$ & $(64, 32)$: Wrong.
- Bard: $[((8.33, 5), (8.33, 22.5)), ((8.33, 22.5), (25, 22.5)), ((8.33, 22.5), (0, 25))]$: Wrong.
- Hugging Chat: $[((12.5, 0), (12.5, 22.5)), ((12.5, 22.5), (25.0, 22.5)), ((25.0, 22.5), (25.0, 25.0))]$: a bit wrong on the last line
- GPT-4:

$[((12.5, 0), (12.5, 22.5)), ((12.5, 22.5), (0, 22.5)), ((12.5, 22.5), (25.0, 22.5))]$ just change in sequence of lines, and that's inconsequential, so the answer is correct.

4. Conclusion

Traditional methods of computing midcurves are predominantly rules-based and thus, have limitation of not developing a generic model which will accept any input shape. This paper presents a novel Large Language Model based approach which attempts to build such a generic model. One such model, GPT-4, seems to be very effective. Although other proprietary and open-source models need to catch-up with GPT-4, even GPT-4 needs to be developed further to understand not just sequential lines but graphs/networks with different shapes, essentially, the geometry.

REFERENCES

- MidcurveNN: Encoder-Decoder Neural Network for Computing Midcurve of a Thin Polygon, viXra.org e-Print archive, viXra:1904.0429 <http://vixra.org/abs/1904.0429>
- CAD Conference 2021, Barcelona, pages 223-225 http://www.cad-conference.net/files/CAD21/CAD21_223-225.pdf
- CAD & Applications 2022 Journal paper 19(6) [http://www.cad-journal.net/files/vol_19/CAD_19\(6\)_2022_1154-1161.pdf](http://www.cad-journal.net/files/vol_19/CAD_19(6)_2022_1154-1161.pdf)